

Simple Linear Regression an Logistic Regression

EBRU KILIÇ

Last compiled on 5 November 2022

Contents

1 Part 1: Simple Linear Regression	1
2 Part 1: Solution	1
3 Part 2: Logistic Regression	8
4 Part 2: Solution	9
5 References	12

1 Part 1: Simple Linear Regression

Exam grades and weekly spent time on self study x (in hours) of 14 statistics students are given in the following table.

Self study	25.0	26.2	24.9	23.7	22.8	24.6	23.6	23.0	22.5	26.2	25.8	24.0	22.1	21.7
Exam Grades	63	53	52	46	34	47	43	37	40	45	53	42	32	49

1. Create a data frame in R with the above data. Plot the data with the weekly spent time on self study in the x -axis and exam grades on the y -axis (You should include labels for your axes and a title for the plot)
2. Obtain the least squares regression line of exam grades on weekly spent time on self study. Interpret your model result (Using the whole data set)
3. Fit the linear model after partitioning your data set into training and testing (round the number of observations when it is necessary). After fitting the model, compare your parameter estimates with the model result in Question 2. Then, make predictions on testing data and compare with the original observations.
4. Using the `plot` command, comment on the validity of the assumption of the model that you fit in Question 3 (Note before using the `plot` command you may wish to specify a 2x2 graphics window using `par(mfrow = c(2, 2))`).
5. Calculate a 95% confidence interval for the slope regression parameter for the last model you fit in Question 3. (Note that the number of degrees of freedom should be obtained from the R output). For this you can use a built-in function in R

2 Part 1: Solution

```
#install.packages("modelr")
#install.packages("tinytex")
#tinytex::install_tinytex()
#install.packages("mlbench")
#install.packages("caret")
#install.packages("ggplot2")
#install.packages("lattice")
```

```
# Last 5 digits of "14449202680"
# and par() method are used because of the reproducibility and better plots.
```

```
set.seed(02680)
par(mfrow = c(2, 3))
```

```
# Data set on the given pdf file
datas <- data.frame(
  study = c(25.0, 26.2, 24.9, 23.7, 22.8, 24.6, 23.6, 23.0, 22.5, 26.2, 25.8, 24.0, 22.1, 21.7),
  grades = c(63, 53, 52, 46, 34, 47, 43, 37, 40, 45, 53, 42, 32, 49))
```

```
# All data's regression
allreggrades <- lm(grades ~ study, data = datas)
```


```
# Control the outliers. There is no problem since it is not outlier.
z_score <- as.data.frame(sapply(datas, function(datas)
  (abs(datas-mean(datas))/sd(datas))))
dim(datas)
```

```
## [1] 14 2
```

```
except_outliers <- z_score[!rowSums(3 < z_score), ]
dim(except_outliers)
```

```
## [1] 14 2
```

```
# Plot of Study and grades
# least squares regression line
# plot() and abline() has used
plot(datas)
abline(allreggrades)
```



```
# Data partitioning as %80.
```

```
index <- sample(seq_len(nrow(datas)), round(8/10*nrow(datas)))
length(index)
```

```
## [1] 11
```

```
# Training data
```

```
train = datas[index, ]
```

```
# Test data
```

```
test <- datas[-index, ]
```

```
### FITTING THE MODEL
```

```
regressionDatas <- lm(grades ~ study, data = train)
```

```
# the P-Value < 0.05
```

```
# it is more meaningful since it has got lower predictor value then p-value
```

```
summary(regressionDatas)
```

```
##
```

```
## Call:
```

```
## lm(formula = grades ~ study, data = train)
```

```
##
```

```
## Residuals:
```

```
##      Min       1Q   Median       3Q      Max
## -8.327  -4.759  -1.499   1.790  14.600
```


```
##
```



```
## Coefficients:
```


```
##              Estimate Std. Error t value Pr(>|t|)
```

```
## (Intercept)  -20.618    39.587  -0.521    0.615
## study        2.761     1.654   1.669    0.129
##
## Residual standard error: 7.314 on 9 degrees of freedom
## Multiple R-squared:  0.2364, Adjusted R-squared:  0.1516
## F-statistic: 2.787 on 1 and 9 DF,  p-value: 0.1294
```

```
plot(regressionDatas)
```







```
# confidence interval = 95%
# slope regression parameter for the model.
confint(regressionDatas, level = 95/100)
```

```
##              2.5 %   97.5 %
## (Intercept) -110.1694047 68.93349
## study       -0.9803677  6.50184
```


```
# Shows the predictions on testing data.
# It shows the differences between predicted values
# While we are making prediction
predict(regressionDatas, train)
```

```
##      8      1      6      14      5      12      11      4
## 42.87897 48.40045 47.29615 39.29002 42.32683 45.63971 50.60904 44.81149
##     10      9      7
## 51.71333 41.49861 44.53542
```

```
train$grades
```

```
## [1] 37 63 47 49 34 42 53 46 45 40 43
```

```
# Shows the residuals that coming from linear regression.
plot(predict(regressionDatas), rstudent(regressionDatas))
```



```
# Model's performance
```

```
library(modelr)
```

```
summary(regressionDatas)
```

```
##
```

```
## Call:
```

```
## lm(formula = grades ~ study, data = train)
```

```
##
```

```
## Residuals:
```

```
##      Min       1Q   Median       3Q      Max
## -8.327 -4.759 -1.499   1.790 14.600
```

```
##
```

```
## Coefficients:
```

```
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  -20.618     39.587  -0.521   0.615
## study           2.761       1.654   1.669   0.129
```

```
##
```

```
## Residual standard error: 7.314 on 9 degrees of freedom
```

```
## Multiple R-squared:  0.2364, Adjusted R-squared:  0.1516
```

```
## F-statistic: 2.787 on 1 and 9 DF,  p-value: 0.1294
```

```
data.frame(
```

```
  #Root Mean Square Error
```

```
  RMSE = rmse(regressionDatas, data = train),
```

```
  #Mean Absolute Error
```

```
  MAE = mae(regressionDatas, data = train), AIC(regressionDatas),
```

```
  BIC(regressionDatas)
```

```
)
```

```
##          RMSE          MAE AIC.regressionDats. BIC.regressionDats.
## 1 6.61586 5.07073          78.78498          79.97867
```

```
# Relationship's strength between the dependent variable on a convenient
# and my model is measuring by r-squared.
```

3 Part 2: Logistic Regression

Consider the available example data set below

```
# Last 5 digits of "14449202680"
# and par() method are used because of the reproducibility and better plots.
set.seed(02680)
par(mfrow=c(2, 3))
```

```
# install.packages("mlbench")
library(mlbench)
data(BreastCancer)
summary(BreastCancer)
```

```
##          Id          Cl.thickness  Cell.size  Cell.shape  Marg.adhesion
## Length:699          1          :145  1          :384  1          :353  1          :407
## Class :character    5          :130  10         : 67  2          : 59  2          : 58
## Mode  :character    3          :108  3          : 52  10         : 58  3          : 58
##          4          : 80  2          : 45  3          : 56  10         : 55
##          10         : 69  4          : 40  4          : 44  4          : 33
##          2          : 50  5          : 30  5          : 34  8          : 25
##          (Other):117  (Other): 81  (Other): 95  (Other): 63
## Epith.c.size  Bare.nuclei  Bl.cromatin  Normal.nucleoli  Mitoses
## 2          :386  1          :402  2          :166  1          :443  1          :579
## 3          : 72  10         :132  3          :165  10         : 61  2          : 35
## 4          : 48  2          : 30  1          :152  3          : 44  3          : 33
## 1          : 47  5          : 30  7          : 73  2          : 36  10         : 14
## 6          : 41  3          : 28  4          : 40  8          : 24  4          : 12
## 5          : 39  (Other): 61  5          : 34  6          : 22  7          : 9
## (Other): 66  NA's : 16  (Other): 69  (Other): 69  (Other): 17
##          Class
## benign :458
## malignant:241
##
##
##
##
##
```

```
# You can check the details here
# https://www.rdocumentation.org/packages/mlbench/versions/2.1-3/topics/BreastCancer
```


1. Convert your Class variable into a numerical one since you have two classes (benign malignant) you can make it one of them as 0 and the other one is 1
2. Fit a logistic regression model to classify **Class** using Mitoses (DO NOT FORGET TO PARTITION YOUR DATA INTO TRAINING AND TESTING DATA SETS, DO NOT FORGET THAT THIS DATA SET INCLUDES QUALITATIVE PREDICTORS !)
3. Make predictions and compare with the true observations (using TEST DATA SET). Calculate and interpret the Confusion Matrix results
4. Fit a multiple logistic regression to classify **Class** by using more than one predictor
5. Compare simple logistic and multiple logistic regression models using F1-score to make a decision on the best model. Why the overall accuracy is not enough as a performance measure ? Explain shortly

4 Part 2: Solution

Use the given R-code chunk below to make your calculations and summarize your result thereafter by adding comments on it,

- MAKE SURE THAT ALL NECESSARY PACKAGES ARE ALREADY INSTALLED and READY TO USE
- You can use as many as Rcode chunks you want. In the final output, both Rcodes and your outputs including your comments should appear in an order

```
# You can check the details here
#      https://www.rdocumentation.org/packages/mlbench/versions/2.1-3/topics/BreastCancer

# numerical variable has comprised from class variable
# "malignant" = 1 , "benign" = 0

vectorIndex = c()
for(m in 1:length(BreastCancer$Class))
{
  if("malignant" == BreastCancer$Class[m])
  {
    vectorIndex[m] = 1;
  }
  else if ("benign" == BreastCancer$Class[m])
  {
    vectorIndex[m] = 0;
  }
}

# NewBreastCancer is the new fixed data
NewBreastCancer = BreastCancer
NewBreastCancer$Class = vectorIndex

# Data partitioning as %80.
traIndex <- sample(seq_len(nrow(NewBreastCancer)),
                   round(8/10*nrow(NewBreastCancer)))
length(traIndex)
```

```
## [1] 559
```

```
# BreastCancer's train data
```

```
TrainBreastCancer = NewBreastCancer[traIndex, ]
```

```
# BreastCancer's test data
```

```
TestBreastCancer <- NewBreastCancer[-traIndex, ]
```

```
# model fitting
```

```
FitBreastCancer = glm(Class ~ Mitoses, data = TrainBreastCancer,  
                      family = "binomial")
```

```
summary(FitBreastCancer)
```

```
##
```

```
## Call:
```

```
## glm(formula = Class ~ Mitoses, family = "binomial", data = TrainBreastCancer)
```

```
##
```

```
## Deviance Residuals:
```

```
##      Min       1Q   Median       3Q      Max  
## -2.2649  -0.7491  -0.7491   0.4001   1.6781
```

```
##
```

```
## Coefficients:
```

```
##              Estimate Std. Error z value Pr(>|z|)  
## (Intercept)  -1.1274     0.1078 -10.462  < 2e-16 ***  
## Mitoses2       2.4267     0.4730   5.130 2.89e-07 ***  
## Mitoses3       3.6123     0.7438   4.856 1.20e-06 ***  
## Mitoses4      18.6935    1398.7210   0.013   0.989  
## Mitoses5      18.6935    1769.2576   0.011   0.992  
## Mitoses6      18.6935    2284.1018   0.008   0.993  
## Mitoses7      18.6935    1615.1039   0.012   0.991  
## Mitoses8      18.6935    1615.1039   0.012   0.991  
## Mitoses10     18.6935    1192.8333   0.016   0.987
```

```
## ---
```

```
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
##
```

```
## (Dispersion parameter for binomial family taken to be 1)
```

```
##
```

```
## Null deviance: 727.90 on 558 degrees of freedom
```

```
## Residual deviance: 561.73 on 550 degrees of freedom
```

```
## AIC: 579.73
```

```
##
```

```
## Number of Fisher Scoring iterations: 16
```

```
# Fitting model's Prediction part
```

```
FitPreBreastCancer = ifelse(predict(FitBreastCancer, type = "response")  
                             > 0.5, "1", "0")
```

```
# comparison
```

```
table(TrainBreastCancer$Class)
```

```
##
```

```
## 0 1
```

```
## 360 199
```

```
table(FitPreBreastCancer)
```

```
## FitPreBreastCancer
##    0    1
## 466  93
```

```
# error rate of training
```

```
# mean error rate
```

```
error = function(actual, predicted)
{
  mean(predicted != actual)
}
```

```
error(actual = TrainBreastCancer$Class, predicted = FitPreBreastCancer)
```

```
## [1] 0.2182469
```

```
FitPreBreastCancer = ifelse(predict(FitBreastCancer, type = "response",
                                   newdata = TestBreastCancer)
                             > 0.5, "1", "0")
error(actual = TestBreastCancer$Class, predicted = FitPreBreastCancer)
```

```
## [1] 0.1785714
```

```
# Calculate the Confusion Matrix on training data
```

```
library(caret)
```

```
## Loading required package: ggplot2
```

```
## Loading required package: lattice
```

```
# Calculating and interpreting the Confusion Matrix results.
```

```
confusionMatrix(as.factor(FitPreBreastCancer),
                 as.factor(TestBreastCancer$Class), mode = "everything")
```

```
## Confusion Matrix and Statistics
```

```
##
```

```
##           Reference
```

```
## Prediction 0  1
```

```
##           0 93 20
```

```
##           1  5 22
```

```
##
```

```
##           Accuracy : 0.8214
```

```
##           95% CI : (0.7478, 0.881)
```

```
## No Information Rate : 0.7
```

```
## P-Value [Acc > NIR] : 0.0007454
```

```
##
```

```
##           Kappa : 0.5265
```

```
##
```

```
## McNemar's Test P-Value : 0.0051103
```

```
##
##      Sensitivity : 0.9490
##      Specificity : 0.5238
##      Pos Pred Value : 0.8230
##      Neg Pred Value : 0.8148
##      Precision : 0.8230
##      Recall : 0.9490
##      F1 : 0.8815
##      Prevalence : 0.7000
##      Detection Rate : 0.6643
##      Detection Prevalence : 0.8071
##      Balanced Accuracy : 0.7364
##
##      ' Positive' Class : 0
##
```

5 References

<https://www.r-project.org/> <https://cran.rstudio.com/> <https://www.reddit.com/r/programming/> <https://www.w3schools.com/r/>